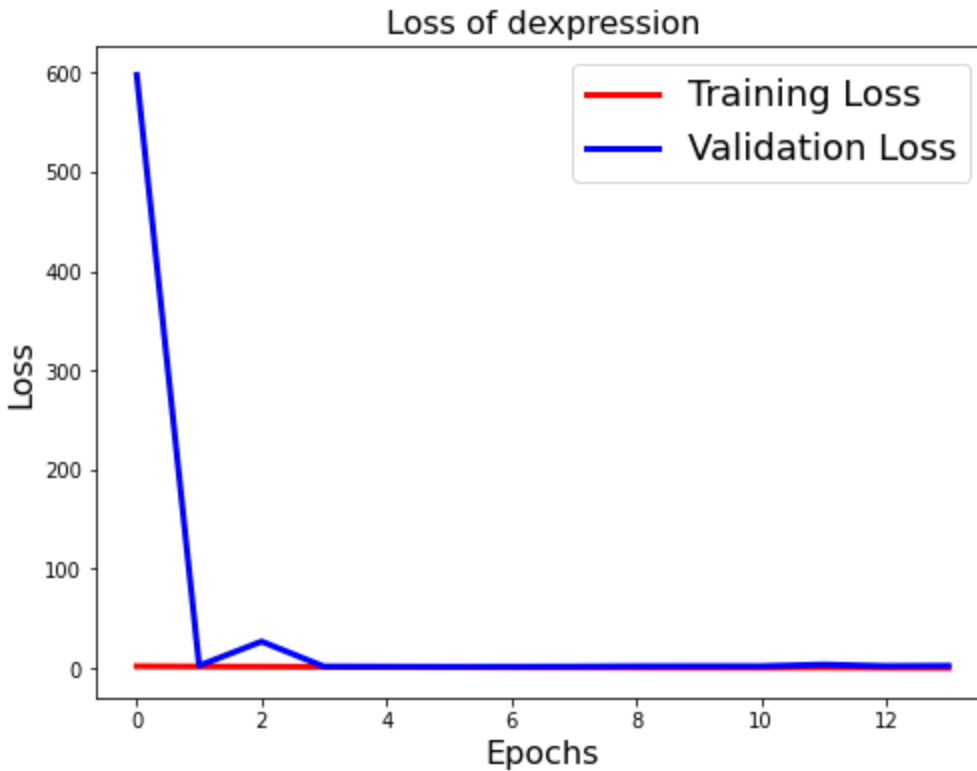

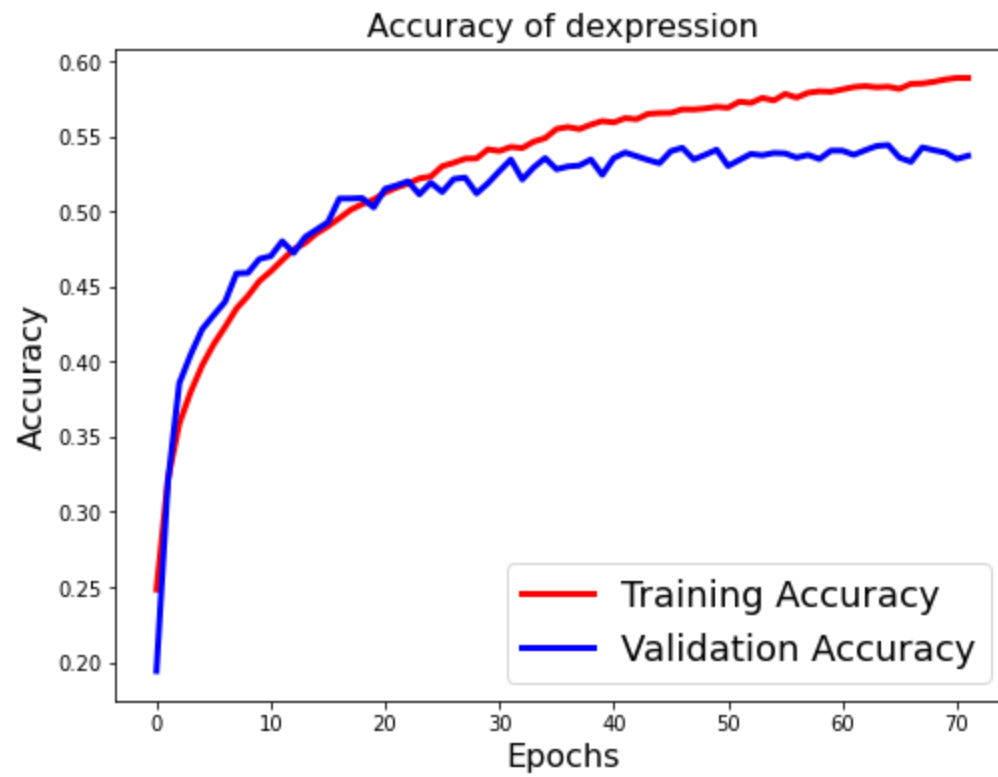



- **Implementation of resNet with data augmentation**



Testing Loss = 2.389908790588379
Testing Accuracy = 0.49623849987983704

- **Implementation of CNN_1 with data Augmentation**







Testing Loss = 1.2271398305892944
Testing Accuracy = 0.5416550636291504

loss: 1.1023 - accuracy: 0.5827 - val_loss: 1.2072 - val_accuracy: 0.5437 - lr: 6.5610e-04

- **Tiny_Xception with data augmentation**

Testing Loss = 1.25651216506958
Testing Accuracy = 0.5285595059394836



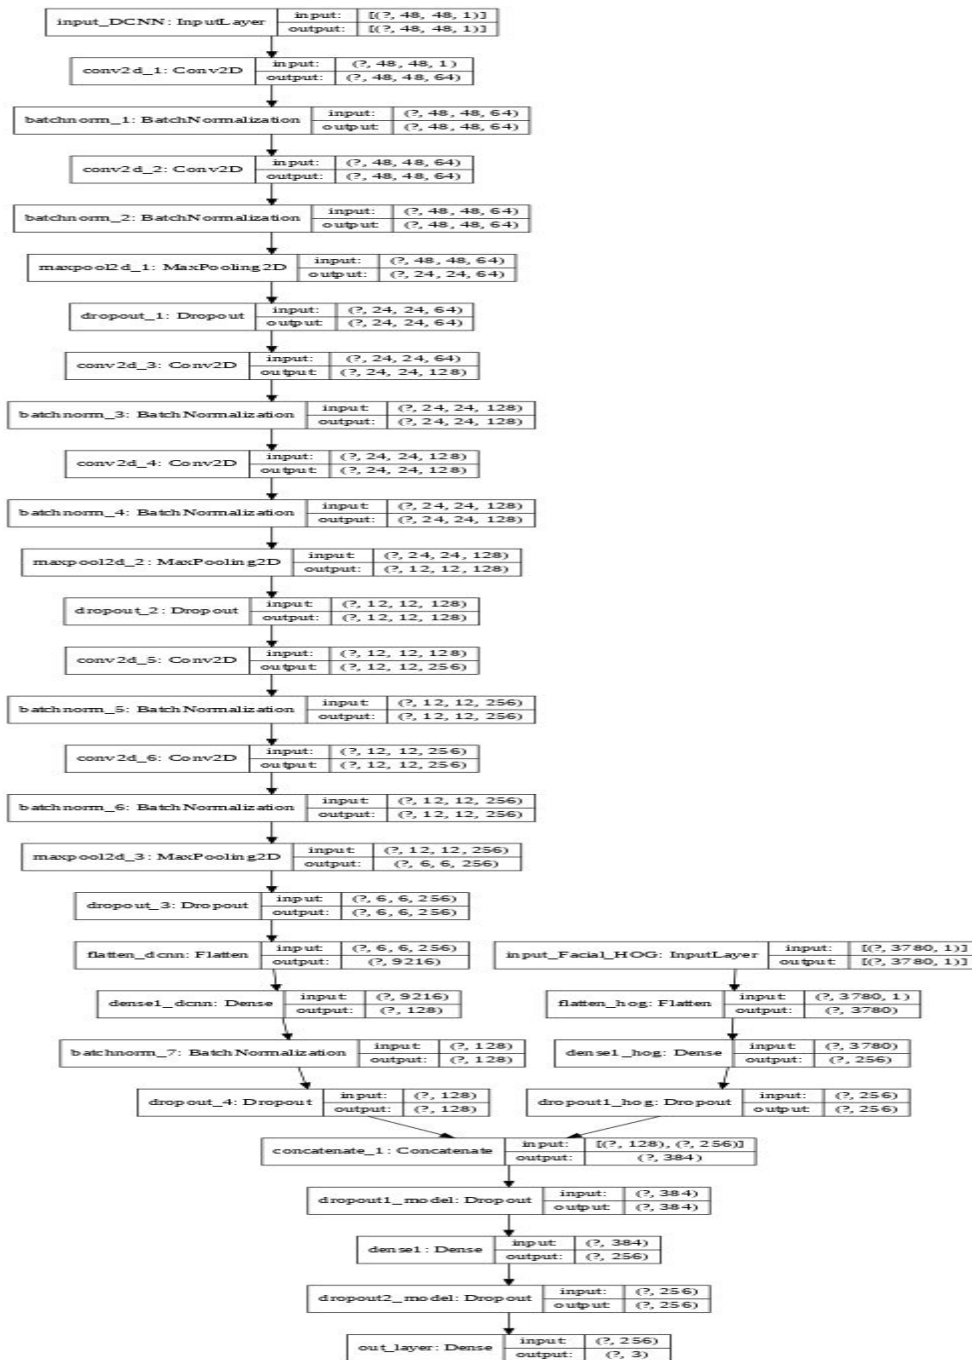
loss: 1.0381 - accuracy: 0.6147 - val_loss: 1.2303 - val_accuracy: 0.5344 - lr: 4.7830e-04

• Start with HOG

- 3 labels \Rightarrow accuracy 80%
- 7 labels \Rightarrow accuracy 65%
- 7 labels with data augmentation \Rightarrow 69%
- 7 labels with data balancing \Rightarrow 73%


1. 3 labels (sad, happy and neutral)


Model :



Train :

```
Epoch 31/100
599/598 [=====] - 12s 20ms/step - loss: 0.4873 - accuracy: 0.7987 - val_loss: 0.4841 - val_accuracy: 0.8072 - lr:
Epoch 32/100
599/598 [=====] - 12s 20ms/step - loss: 0.4845 - accuracy: 0.7991 - val_loss: 0.4840 - val_accuracy: 0.8025 - lr:
Epoch 33/100
599/598 [=====] - 12s 20ms/step - loss: 0.4854 - accuracy: 0.8007 - val_loss: 0.4785 - val_accuracy: 0.8039 - lr:
Epoch 34/100
597/598 [=====>.] - ETA: 0s - loss: 0.4786 - accuracy: 0.7996
Epoch 00034: ReduceLRonPlateau reducing learning rate to 6.25000029685907e-05.
599/598 [=====] - 12s 20ms/step - loss: 0.4793 - accuracy: 0.7994 - val_loss: 0.4886 - val_accuracy: 0.7969 - lr:
Epoch 35/100
599/598 [=====] - 12s 20ms/step - loss: 0.4663 - accuracy: 0.8060 - val_loss: 0.4754 - val_accuracy: 0.8035 - lr:
Epoch 36/100
598/598 [=====>.] - ETA: 0s - loss: 0.4607 - accuracy: 0.8109Restoring model weights from the end of the best epoch
599/598 [=====] - 12s 20ms/step - loss: 0.4607 - accuracy: 0.8108 - val_loss: 0.4770 - val_accuracy: 0.8054 - lr:
Epoch 00036: early stopping
```





Accuracy :


total wrong validation predictions: 410


	precision	recall	f1-score	support
0	0.92	0.92	0.92	899
1	0.76	0.72	0.74	608
2	0.70	0.74	0.72	620
accuracy			0.81	2127
macro avg	0.79	0.79	0.79	2127
weighted avg	0.81	0.81	0.81	2127

2. 7 labels:

Train:

Epoch 34/100
 1010/1009 [=====] - 15s 14ms/step - loss: 0.9711 - accuracy: 0.6402 - val_loss: 0.9282 - val_accuracy: 0.6556 - l
 Epoch 35/100
 1010/1009 [=====] - 15s 15ms/step - loss: 0.9673 - accuracy: 0.6426 - val_loss: 0.9273 - val_accuracy: 0.6514 - l
 Epoch 36/100
 1010/1009 [=====] - 15s 14ms/step - loss: 0.9599 - accuracy: 0.6439 - val_loss: 0.9287 - val_accuracy: 0.6478 - l
 Epoch 37/100
 1009/1009 [=====>.] - ETA: 0s - loss: 0.9599 - accuracy: 0.6469
 Epoch 00037: ReduceLROnPlateau reducing learning rate to 6.25000029685907e-05.
 1010/1009 [=====] - 15s 15ms/step - loss: 0.9600 - accuracy: 0.6469 - val_loss: 0.9222 - val_accuracy: 0.6548 - l
 Epoch 38/100
 1010/1009 [=====] - 14s 14ms/step - loss: 0.9487 - accuracy: 0.6507 - val_loss: 0.9200 - val_accuracy: 0.6539 - l
 Epoch 39/100
 1009/1009 [=====>.] - ETA: 0s - loss: 0.9458 - accuracy: 0.6509
 Restoring model weights from the end of the best epoch
 1010/1009 [=====] - 15s 14ms/step - loss: 0.9457 - accuracy: 0.6510 - val_loss: 0.9203 - val_accuracy: 0.6556 - l
 Epoch 00039: early stopping





Accuracy :

total wrong validation predictions: 1236

	precision	recall	f1-score	support
0	0.55	0.68	0.61	495
1	0.83	0.09	0.16	55
2	0.64	0.27	0.38	512
3	0.87	0.87	0.87	899
4	0.55	0.48	0.51	608
5	0.72	0.81	0.76	400
6	0.54	0.76	0.63	620
accuracy			0.66	3589
macro avg	0.67	0.57	0.56	3589
weighted avg	0.67	0.66	0.64	3589


3. 7 labels + data balancing :


Train:

```

1231/1230 [=====] - 17s 14ms/step - loss: 0.9526 - accuracy: 0.6446 - val_loss: 0.8738 - val_accuracy: 0.6616 -
Epoch 36/100
1231/1230 [=====] - 18s 14ms/step - loss: 0.9455 - accuracy: 0.6448 - val_loss: 0.8732 - val_accuracy: 0.6658 -
Epoch 37/100
1231/1230 [=====] - 17s 14ms/step - loss: 0.9490 - accuracy: 0.6437 - val_loss: 0.8758 - val_accuracy: 0.6610 -
Epoch 38/100
1231/1230 [=====] - 18s 14ms/step - loss: 0.9514 - accuracy: 0.6440 - val_loss: 0.8718 - val_accuracy: 0.6603 -
Epoch 39/100
1230/1230 [=====>.] - ETA: 0s - loss: 0.9449 - accuracy: 0.6461
Epoch 00039: ReduceLROnPlateau reducing learning rate to 1.5625000742147677e-05.
1231/1230 [=====] - 17s 14ms/step - loss: 0.9450 - accuracy: 0.6461 - val_loss: 0.8778 - val_accuracy: 0.6566 -
Epoch 40/100
1231/1230 [=====] - 17s 14ms/step - loss: 0.9407 - accuracy: 0.6458 - val_loss: 0.8706 - val_accuracy: 0.6623 -
Epoch 41/100
1230/1230 [=====>.] - ETA: 0s - loss: 0.9370 - accuracy: 0.6490Restoring model weights from the end of the best e
1231/1230 [=====] - 17s 14ms/step - loss: 0.9371 - accuracy: 0.6490 - val_loss: 0.8687 - val_accuracy: 0.6628 -
Epoch 00041: early stopping

```





Accuracy:


total wrong validation predictions: 1463


	precision	recall	f1-score	support
0	0.52	0.62	0.56	592
1	0.85	0.86	0.86	438
2	0.54	0.17	0.26	616
3	0.89	0.89	0.89	899
4	0.48	0.56	0.51	608
5	0.76	0.79	0.77	604
6	0.57	0.73	0.64	620
accuracy			0.67	4377
macro avg	0.66	0.66	0.64	4377
weighted avg	0.67	0.67	0.65	4377

4. Training more with 7 labels with data balancing :

Train:

Epoch 31/100
 1231/1230 [=====] - 17s 14ms/step - loss: 0.8452 - accuracy: 0.6858 - val_loss: 0.8162 - val_accuracy: 0.6916 - l
 Epoch 32/100
 1230/1230 [=====>.] - ETA: 0s - loss: 0.8443 - accuracy: 0.6875
 Epoch 00032: ReduceLRonPlateau reducing learning rate to 1.5625000742147677e-05.
 1231/1230 [=====] - 17s 14ms/step - loss: 0.8443 - accuracy: 0.6875 - val_loss: 0.8084 - val_accuracy: 0.6929 - l
 Epoch 33/100
 1231/1230 [=====] - 17s 14ms/step - loss: 0.8414 - accuracy: 0.6895 - val_loss: 0.8101 - val_accuracy: 0.6916 - l
 Epoch 34/100
 1231/1230 [=====] - 17s 14ms/step - loss: 0.8386 - accuracy: 0.6909 - val_loss: 0.8112 - val_accuracy: 0.6925 - l
 Epoch 35/100
 1230/1230 [=====>.] - ETA: 0s - loss: 0.8337 - accuracy: 0.6921Restoring model weights from the end of the best epc
 Epoch 00035: ReduceLRonPlateau reducing learning rate to 3.906250185536919e-06.
 1231/1230 [=====] - 18s 14ms/step - loss: 0.8337 - accuracy: 0.6921 - val_loss: 0.8115 - val_accuracy: 0.6934 - l
 Epoch 00035: early stopping





Accuracy:

total wrong validation predictions: 1340

	precision	recall	f1-score	support
0	0.57	0.60	0.59	592
1	0.89	0.92	0.90	438
2	0.59	0.26	0.36	616
3	0.90	0.89	0.90	899
4	0.52	0.57	0.54	608
5	0.74	0.84	0.79	604
6	0.58	0.75	0.65	620
accuracy			0.69	4377
macro avg	0.69	0.69	0.68	4377
weighted avg	0.69	0.69	0.68	4377

5. With more training :

Train:

Epoch 15/100

1230/1230 [=====>.] - ETA: 0s - loss: 0.8107 - accuracy: 0.6997

Epoch 00015: ReduceLRonPlateau reducing learning rate to 1.5625000742147677e-05.

1231/1230 [=====] - 17s 14ms/step - loss: 0.8108 - accuracy: 0.6996 - val_loss: 0.8037 - val_accuracy: 0.6941 - l

Epoch 16/100


1231/1230 [=====] - 17s 14ms/step - loss: 0.8123 - accuracy: 0.6969 - val_loss: 0.8017 - val_accuracy: 0.6955 - l


Epoch 17/100

1229/1230 [=====>.] - ETA: 0s - loss: 0.8091 - accuracy: 0.7007Restoring model weights from the end of the best epo

1231/1230 [=====] - 17s 14ms/step - loss: 0.8089 - accuracy: 0.7009 - val_loss: 0.7989 - val_accuracy: 0.6959 - l

Epoch 00017: early stopping





Accuracy:

› total wrong validation predictions: 1331

	precision	recall	f1-score	support
0	0.58	0.65	0.61	592
1	0.88	0.93	0.90	438
2	0.59	0.28	0.38	616
3	0.89	0.89	0.89	899
4	0.53	0.57	0.55	608
5	0.75	0.82	0.78	604
6	0.58	0.72	0.64	620
accuracy			0.70	4377
macro avg	0.69	0.69	0.68	4377
weighted avg	0.69	0.70	0.69	4377

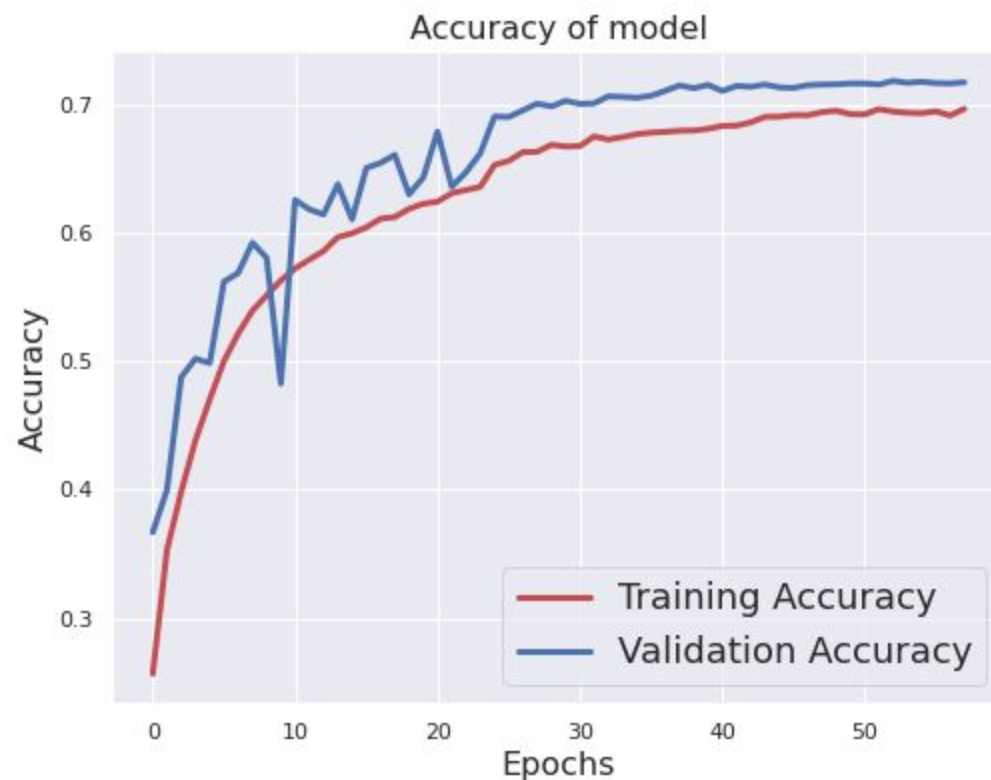
6. 7 labels with more data balancing with disgust to 5470 :


Train:

```

Epoch 54/100
1291/1290 [=====] - 18s 14ms/step - loss: 0.8243 - accuracy: 0.6934 - val_loss: 0.7692 - val_accuracy: 0.7167 - l
Epoch 55/100
1291/1290 [=====] - 18s 14ms/step - loss: 0.8219 - accuracy: 0.6930 - val_loss: 0.7654 - val_accuracy: 0.7175 - l
Epoch 56/100
1290/1290 [=====>.] - ETA: 0s - loss: 0.8212 - accuracy: 0.6945
Epoch 00056: ReduceLROnPlateau reducing learning rate to 3.906250185536919e-06.
1291/1290 [=====] - 18s 14ms/step - loss: 0.8211 - accuracy: 0.6945 - val_loss: 0.7714 - val_accuracy: 0.7164 - l
Epoch 57/100
1291/1290 [=====] - 18s 14ms/step - loss: 0.8271 - accuracy: 0.6912 - val_loss: 0.7705 - val_accuracy: 0.7162 - l
Epoch 58/100
1288/1290 [=====>.] - ETA: 0s - loss: 0.8204 - accuracy: 0.6964Restoring model weights from the end of the best epc
1291/1290 [=====] - 18s 14ms/step - loss: 0.8201 - accuracy: 0.6965 - val_loss: 0.7707 - val_accuracy: 0.7171 - l
Epoch 00058: early stopping

```





Accuracy:

total wrong validation predictions: 1293


	precision	recall	f1-score	support
0	0.60	0.64	0.62	591
1	0.89	0.98	0.93	657
2	0.64	0.28	0.39	616
3	0.88	0.92	0.90	899
4	0.54	0.57	0.55	608
5	0.80	0.79	0.80	597
6	0.58	0.74	0.65	620
accuracy			0.72	4588
macro avg	0.70	0.70	0.69	4588
weighted avg	0.72	0.72	0.71	4588

7. More train:

Train:

Epoch 16/100
 1291/1290 [=====] - 18s 14ms/step - loss: 0.7742 - accuracy: 0.7142 - val_loss: 0.7491 - val_accuracy: 0.7299 - lr
 Epoch 17/100
 1291/1290 [=====] - 18s 14ms/step - loss: 0.7717 - accuracy: 0.7134 - val_loss: 0.7457 - val_accuracy: 0.7269 - lr
 Epoch 18/100
 1291/1290 [=====] - 18s 14ms/step - loss: 0.7687 - accuracy: 0.7145 - val_loss: 0.7498 - val_accuracy: 0.7293 - lr
 Epoch 19/100
 1290/1290 [=====>.] - ETA: 0s - loss: 0.7668 - accuracy: 0.7151
 Epoch 00019: ReduceLRonPlateau reducing learning rate to 1.5625000742147677e-05.
 1291/1290 [=====] - 18s 14ms/step - loss: 0.7670 - accuracy: 0.7151 - val_loss: 0.7438 - val_accuracy: 0.7284 - lr
 Epoch 20/100
 1291/1290 [=====] - 18s 14ms/step - loss: 0.7635 - accuracy: 0.7184 - val_loss: 0.7505 - val_accuracy: 0.7265 - lr
 Epoch 21/100
 1288/1290 [=====>.] - ETA: 0s - loss: 0.7596 - accuracy: 0.7202Restoring model weights from the end of the best epoc
 1291/1290 [=====] - 18s 14ms/step - loss: 0.7596 - accuracy: 0.7202 - val_loss: 0.7459 - val_accuracy: 0.7289 - lr
 Epoch 00021: early stopping





Accuracy:

total wrong validation predictions: 1239

	precision	recall	f1-score
0	0.63	0.66	0.64
1	0.89	0.99	0.94
2	0.65	0.33	0.44
3	0.88	0.91	0.89
4	0.57	0.54	0.56
5	0.78	0.82	0.80
6	0.60	0.76	0.67
accuracy	0.73	0.73	0.73

total wrong validation predictions: 1239

	precision	recall	f1-score	support
0	0.63	0.66	0.64	591
1	0.89	0.99	0.94	657
2	0.65	0.33	0.44	616
3	0.88	0.91	0.89	899
4	0.57	0.54	0.56	608
5	0.78	0.82	0.80	597
6	0.60	0.76	0.67	620
accuracy			0.73	4588
macro avg	0.71	0.72	0.70	4588
weighted avg	0.73	0.73	0.72	4588

- **Simple CNN_2**

train:

Epoch 49/200

156/156 [=====>.] - ETA: 0s - loss: 0.5006 - accuracy: 0.8181

Epoch 00049: ReduceLRonPlateau reducing learning rate to 1.5625000742147677e-05.

157/156 [=====] - 13s 85ms/step - loss: 0.5006 - accuracy: 0.8180 - val_loss: 0.9084 - val_accuracy: 0.7208 - lr

Epoch 50/200

157/156 [=====] - 13s 85ms/step - loss: 0.4891 - accuracy: 0.8211 - val_loss: 0.9085 - val_accuracy: 0.7219 - lr

Epoch 51/200

157/156 [=====] - 13s 85ms/step - loss: 0.4921 - accuracy: 0.8202 - val_loss: 0.9140 - val_accuracy: 0.7212 - lr

Epoch 52/200

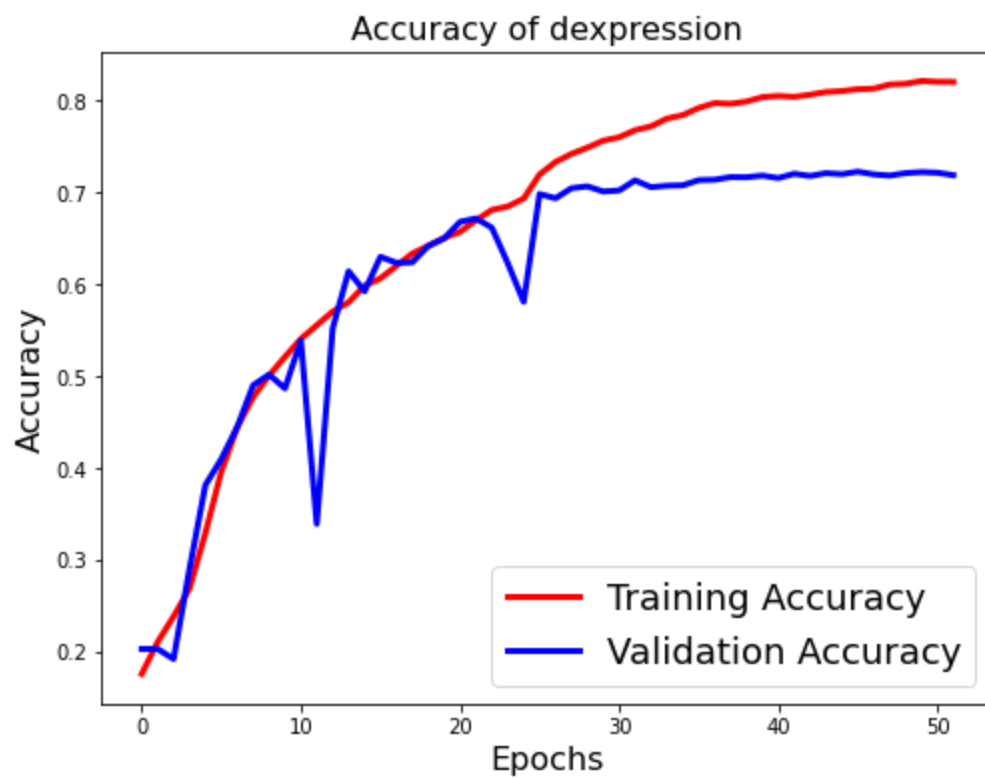
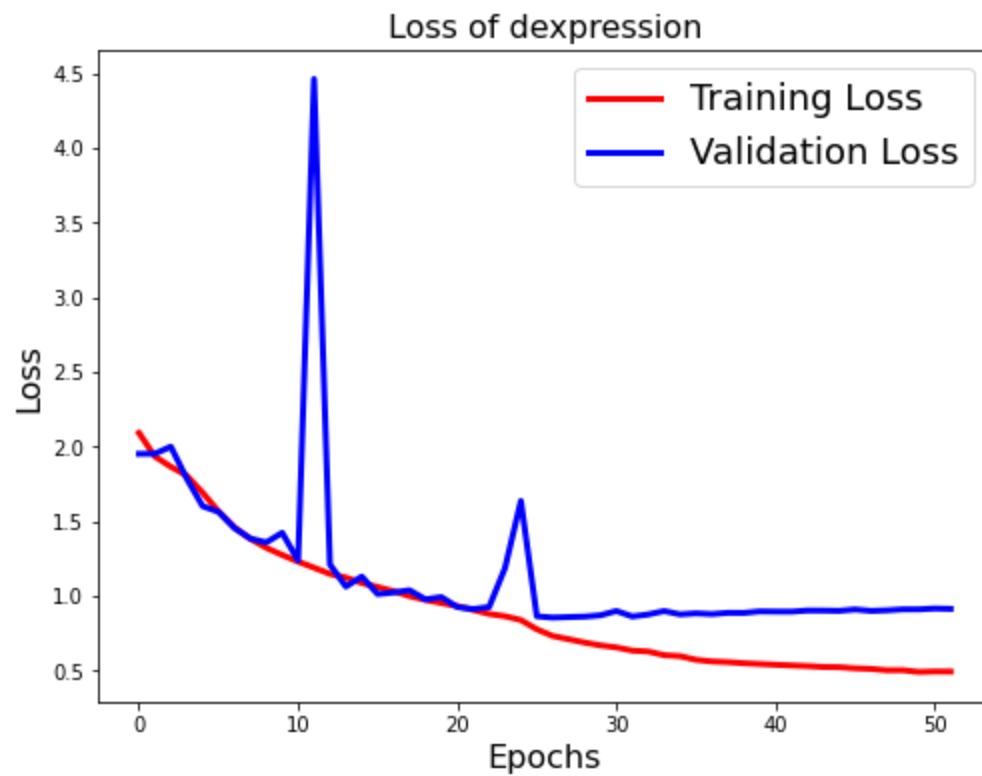
156/156 [=====>.] - ETA: 0s - loss: 0.4917 - accuracy: 0.8200Restoring model weights from the end of the best epoch

Epoch 00052: ReduceLRonPlateau reducing learning rate to 3.906250185536919e-06.

157/156 [=====] - 13s 85ms/step - loss: 0.4914 - accuracy: 0.8200 - val_loss: 0.9115 - val_accuracy: 0.7185 - lr

Epoch 00052: early stopping

Accuracy of degradation



Accuracy:

total wrong validation predictions: 1235

	precision	recall	f1-score	support
0	0.69	0.57	0.63	624
1	0.95	0.97	0.96	559
2	0.60	0.50	0.55	566
3	0.82	0.87	0.84	903
4	0.53	0.60	0.56	587
5	0.83	0.79	0.81	554
6	0.62	0.69	0.65	655
accuracy			0.72	4448
macro avg	0.72	0.71	0.71	4448
weighted avg	0.72	0.72	0.72	4448

- **Ensemble :**

1) Training with 7 models (HoG, resNet and CNNs)

```
[86] from sklearn.metrics import accuracy_score
test_accuracy_score = accuracy_score(rounded_labels, labels_w_average)
print(test_accuracy_score)
```

↳ 0.7052868391451068

```
[87] from sklearn.metrics import accuracy_score
test_accuracy_score = accuracy_score(rounded_labels, labels_average)
print(test_accuracy_score)
```

↳ 0.6674915635545556

```
[88] from sklearn.metrics import accuracy_score
test_accuracy_score = accuracy_score(rounded_labels, labels)
print(test_accuracy_score)
```

↳ 0.49426321709786275

2) Ensemble with two models CNN_2 and dcnn_hog :

```
[135] from sklearn.metrics import accuracy_score  
      test_accuracy_score = accuracy_score(rounded_labels, labels_w_average)  
      print(test_accuracy_score)
```

0.8499437570303712

```
[136] from sklearn.metrics import accuracy_score  
      test_accuracy_score = accuracy_score(rounded_labels, labels_average)  
      print(test_accuracy_score)
```

0.8434195725534308

```
[137] from sklearn.metrics import accuracy_score  
      test_accuracy_score = accuracy_score(rounded_labels, labels)  
      print(test_accuracy_score)
```

0.8137232845894263

DataSet FER2013:



Data Preprocessing

As shown in the charts the data is not balanced, so we augmented the dataset fer2013 to achieve data balancing(Additive Gaussian Noise, Mirror Flipping, Linear Contrast, Sharpen Images and Crop Images).

before	After																																
 <table border="1"> <caption>Emotion Distribution Before Balancing</caption> <thead> <tr> <th>emotion</th> <th>count</th> </tr> </thead> <tbody> <tr><td>0</td><td>5000</td></tr> <tr><td>1</td><td>500</td></tr> <tr><td>2</td><td>5000</td></tr> <tr><td>3</td><td>9000</td></tr> <tr><td>4</td><td>6000</td></tr> <tr><td>5</td><td>4000</td></tr> <tr><td>6</td><td>6000</td></tr> </tbody> </table>	emotion	count	0	5000	1	500	2	5000	3	9000	4	6000	5	4000	6	6000	 <table border="1"> <caption>Emotion Distribution After Balancing</caption> <thead> <tr> <th>emotion</th> <th>count</th> </tr> </thead> <tbody> <tr><td>0</td><td>8500</td></tr> <tr><td>1</td><td>6500</td></tr> <tr><td>2</td><td>9000</td></tr> <tr><td>3</td><td>8500</td></tr> <tr><td>4</td><td>8500</td></tr> <tr><td>5</td><td>8000</td></tr> <tr><td>6</td><td>9000</td></tr> </tbody> </table>	emotion	count	0	8500	1	6500	2	9000	3	8500	4	8500	5	8000	6	9000
emotion	count																																
0	5000																																
1	500																																
2	5000																																
3	9000																																
4	6000																																
5	4000																																
6	6000																																
emotion	count																																
0	8500																																
1	6500																																
2	9000																																
3	8500																																
4	8500																																
5	8000																																
6	9000																																
{0: 'anger', 1: 'disgust', 2: 'fear', 3: 'happiness', 4: 'sadness', 5: 'surprise', 6: 'neutral'}																																	
Dataset size before data balancing = 35887	Dataset size before data balancing = 60061																																

- Results

- References

- [Fer-using-multiple-pipelines](#)
- [Facial expression databases](#)
- [Deep Facial Expression Recognition: A Survey](#)
- [Suppressing Uncertainties for Large-Scale Facial Expression Recognition](#)